



02 Linear Functions



- 01 Vectors
- **02 Linear Functions**
- 03 Norms and Distances
- 04 Clustering
- 05 Linear Independence

Unit 2: Matrices, Book ILA Ch. 6-11 + Book IMC Ch. 2

Unit 3: Least Squares, Book ILA Ch. 12-14 + Book IMC Ch. 8

Unit 4: Eigen-decomposition, Book IMC Ch. 10, 12, 19

Outline: 02 Linear Functions

- [Linear and affine functions](#)
- [Taylor approximation](#)
- [Regression model](#)

Outline: 02 Linear Functions

- [Linear and affine functions](#)
- [Taylor approximation](#)
- [Regression model](#)

Superposition and linear functions

Notation: The notation $f : \mathbb{R}^n \rightarrow \mathbb{R}$ means f is a function mapping n -vectors to numbers.

Definition: We say that f satisfies the superposition property if:

$$f(\alpha x + \beta y) = \alpha f(x) + \beta f(y)$$

holds for all scalars α, β and all n -vectors x, y .

Definition: A function that satisfies superposition is called linear.

Example: the inner product function

Definition: For a an n -vector, the inner product function is defined as

$$f(x) = a^T x = a_1 x_1 + \dots + a_n x_n.$$

We see that $f(x)$ is a weighted sum of the entries of x .

Exercise: Show that the inner product function is linear.

In Python:

```
In [65]: import numpy as np

a = np.array([-2, 0, 1, -3])
f = lambda x: np.inner(a, x); f
```

```
Out[65]: <function __main__.<lambda>(x)>
```

```
In [66]: x, y = np.array([2, 2, -1, 1]), np.array([0, 1, -1, 0])
alpha, beta = 1.5, -3.7
```

```
In [67]: lhs, rhs = f(alpha * x + beta * y), alpha * f(x) + beta * f(y)
lhs, rhs
```

```
Out[67]: (-8.3, -8.3)
```

All linear functions are inner product functions

Proposition:

- If $f: \mathbb{R}^n \rightarrow \mathbb{R}$ is linear,
- Then f can be expressed as $f(x) = a^T x$ for some n -vector a . Specifically, a is given by $a_i = f(e_i)$ for i in $\{1, \dots, n\}$ where e_i is the one-hot vector with entry 1 at i .

Exercise: If $f(x) = a^T x$, show that $a_i = f(e_i)$.

In Python:

```
In [68]: a = np.array([-2, 0, 1, -3])
f = lambda x: np.inner(a, x)

e0 = np.array([1, 0, 0, 0])
f(e0)
```

```
Out[68]: -2
```

Affine functions

Definition: A function that is linear plus a constant is called affine. Its general form is:

$$f(x) = a^T x + b \quad \text{with } a \text{ an } n\text{-vector and } b \text{ a scalar}$$


Proposition: A function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ is affine if and only if:

$$f(\alpha x + \beta y) = \alpha f(x) + \beta f(y)$$

for all scalars α, β with $\alpha + \beta = 1$ and all n -vectors x, y .

Linear versus affine functions

In this plot, f is linear and g is affine, not linear



🧠 Sometimes people refer to affine functions as linear.

Outline: Linear Functions

- [Linear and affine functions](#)
- **Taylor approximation**
- [Regression model](#)

Gradient of a function

Definition: The gradient of f at the n -vector x is defined as the n -vector $\nabla f(x)$ as:


$$\nabla f(x) = \left(\frac{\partial f}{\partial x_1}(x), \dots, \frac{\partial f}{\partial x_n}(x) \right).$$

where $\frac{\partial}{\partial x_i}$ is the partial derivative of f with respect to the component i of x .

Exercise: Consider $f(x) = x_1 + (x_2 - x_1)^3$ defined for any 2-vector x . Compute the gradient of f .

Gradient of a function

Example: Consider the function $f : \mathbb{R}^2 \rightarrow \mathbb{R}$.



The arrows represent the gradient of f at different points in \mathbb{R}^2 .

Gradient of a function

In Python, we can use the package `jax` to compute gradients.

In [69]:

```
import jax

f = lambda x: x[0] + (x[1] - x[0]) ** 3
grad_f = jax.grad(f); grad_f(np.array([1., 2.]))
```

Out[69]: DeviceArray([-2., 3.], dtype=float32)

Remark: There are other packages to automatically compute gradients:

- `autograd`, `pytorch`, `tensorflow`.

DeepMind Releases New JAX Libraries for Neural Networks and Reinforcement Learning



Synced

Follow



Feb 21, 2020 · 3 min read



First-order Taylor approximation


Definition: Take $f : \mathbb{R}^n \rightarrow \mathbb{R}$. The first-order Taylor approximation of f near the n -vector z is

$$\hat{f}(x) = f(z) + \nabla f(z)^T(x - z) = f(z) + \frac{\partial f}{\partial x_1}(z)(x_1 - z_1) + \dots + \frac{\partial f}{\partial x_n}(z)(x_n - z_n).$$

Properties:

- $\hat{f}(z) = f(z)$, i.e. equality at z .
- $\hat{f}(x)$ is very close to $f(z)$ when x is very close to z .

First-order Taylor approximation



- The first-order Taylor approximation provides an affine approximation *near* z to a function that is not affine.
- 🤖 This is a good approximation of f only if x is near z .

First-order Taylor approximation

In Python:

In [70]:

```
f = lambda x: x[0] + (x[1] - x[0]) ** 3
grad_f = jax.grad(f)

z = np.array([1., 2.])
f_hat = lambda x: f(z) + np.inner(grad_f(z), (x - z))
f(z), f_hat(z)
```

Out[70]: (2.0, 2.0)

In [71]:

```
close_to_z = np.array([1.01, 2.02]); print(f(close_to_z), f_hat(close_to_z))
far_from_z = np.array([-1., -2.]); print(f(far_from_z), f_hat(far_from_z))
```

```
2.0403010000000004 2.04
-2.0 -6.0
```

Outline: Linear Functions

- [Linear and affine functions](#)
- [Taylor approximation](#)
- [Regression model](#)

Regression

Definition: A regression model is the affine function of x defined as

$$\hat{y} = x^T w + b$$

where:

- x is an n -vector, called a *feature* or *input* vector,
- the elements x_i 's of x are called *regressors*,
- the n -vector w is called *weight vector*,
- the scalar b is called *offset* or *intercept* or *bias*,
- the scalar \hat{y} is called *prediction* --- of some actual *outcome* y .

Example: House prices

Example: Model of house prices:

- Outcome of interest is y , selling price of house in \$1000, in some location, over some period of time.
- Available regressors are $x = (\text{house area, \# bedrooms})$, where house area is in 1000 sq.ft.
- We are given a regression model with weight vector and offset as:

$$w = (148.73, -18.85), b = 54.40$$

If we are given input regressors for different houses, we can use the regression model to make predictions on their prices.

In Python, we use the regression function to return predictions for each house based on beds and area.

```
In [72]: w, b = np.array([148.73, -18.85]), 54.40
y_hat = lambda x: np.inner(x, w) + b

x_area, x_beds = 0.846, 1
y_hat_price = y_hat(np.array([x_area, x_beds])); y_hat_price
```

```
Out[72]: 161.37557999999999
```

Example: House prices

In Python, we evaluate how our model compares to observed data stored in a csv. We use the package `pandas` to load data from a `.csv`.


```
In [73]: import pandas as pd

data = pd.read_csv("data/02_houses.csv")
```




```
In [74]: y_price = data["price"];
x_area = data["area"]; x_beds = data["beds"]
y_hat_price = np.array([w[0] * x_area + w[1] * x_beds]) + b
```

```
In [75]: fig = plt.figure(figsize=(8, 8)); ax = fig.add_subplot(projection='3d')
ax.scatter(x_area, x_beds, y_price, c=x_beds, s=50);
ax.set_xlabel("Area (1000 sq. feet)", fontsize=20); ax.set_ylabel("Number of Bed
```



We compare the observed prices with our predicted prices.


```
In [76]: fig = plt.figure(figsize=(8, 8)); ax = fig.add_subplot(projection='3d')
ax.scatter(x_area, x_beds, y_price, c=x_beds, s=50);
ax.set_xlabel("Area (1000 sq. feet)", fontsize=20); ax.set_ylabel("Number of Bed
x_beds_grid = np.arange(0, 6, 0.5); x_area_grid = np.arange(0, 4, 0.5)
X, Y = np.meshgrid(x_beds_grid, x_area_grid)
zs = np.array(y_hat(np.array([np.ravel(X), np.ravel(Y)]) .T)); Z = zs.reshape(X.s
ax.plot_surface(X, Y, Z, alpha=0.2, color="blue"); plt.show();
```



In [77]:

```
import matplotlib.pyplot as plt

plt.scatter(y_price, y_hat_price, color='#FEBE11')
plt.plot([0, 800], [0, 800], linestyle='dashed', color='#09847A')
plt.xlabel("Actual price (k$)", fontsize=20)
plt.xlim(0, 800)
plt.ylabel("Predicted price (k$)", fontsize=20)
plt.ylim(0, 800)
plt.grid(True)
```



Conclusion: Linear Functions

- [Linear and affine functions](#)
- [Taylor approximation](#)
- [Regression model](#)

Resources

- Ch. 2 of ILA